

Data Engineering Notes

Created: 2024-10-30

Updated: 2024-11-01

References:

- <https://www.coursera.org/learn/spark-hadoop-snowflake-data-engineering>
- <https://learn.microsoft.com/en-us/training/modules/use-apache-spark-azure-databricks/>
- ChatGPT 4o mini
- Google search

Apache Hadoop: open source ecosystem of software enabling parallel processing of big data.

- Hadoop Distributed File System (HDFS): storage system.
- MapReduce: framework within Hadoop that:
 - maps (distributes) tasks across a cluster of computers that store intermediate results (key-value pair) on disk (slower than in-memory).
 - Intermediate results are grouped by keys and sent to the cluster of computers (parallel processing) responsible for reducing (performing user reduction function) and results are combined.

Apache Spark: built on top of Hadoop but stores intermediate results in-memory instead of on disk.

Leverages parallelism for task completion via horizontal scaling (more nodes).

- Each worker/executor node runs a JVM (multi-threaded). Each has multiple slots based on #cores #cpus of node.
- Parallelized jobs are broken down into stages to be performed in order.

Azure Databricks: parallelized data processing on Apache Spark clusters.

- A notebook instance (*SparkSession* object) controls the driver node which distributes work across worker nodes.

Resilient Distributed Datasets (RDD): immutable fundamental data structure in Spark that facilitates distributed computing.

- Resilient: fault tolerant.
- Distributed: RDDs distributed across nodes in Spark clusters.
- Operations:
 - Transformation (lazy, triggered with action): creates new RDDs after applying transformation e.g., filter.
 - Shuffle is a transformation operation that writes data to disk then re-partitions data. Shuffle defines stage boundaries.
 - Action: return result from computation, e.g., count.
- In pyspark shell (sc: spark context, spark: spark session):
 - `rdd = sc.parallelize(list(range(15)))` # create RDD
 - `rdd2 = rdd.map(lambda x: x*2)` #transformation
 - `rdd3 = rdd.filter(lambda x: x%3==0)` #lazy transformation
 - `rdd3.count()` #action triggers computation

Spark SQL DataFrames: DataSet with named columns.

- Distributed collection of data.
- Abstraction of RDDs.
- Optimized for structured data manipulation.
- Works with (semi)-structured data: CSV, Apache Avro, ORC, Parquet, JSON, relational DBs.
- PySpark integrates with Spark SQL can be used for ETL (Extract, Transform, Load) pipelines.
- Spark SQL is a module within Spark providing SQL capabilities for structured data.

- PySpark is a library allowing Spark to be used in Python.

Big Data Storage Architecture

- **Shared everything**: database installed on the server that runs computations. Efficient for small data.
- **Shared disk**: computation done on separate nodes, but all share one storage. Nodes may need to wait for each other to write to disk, reducing efficiency.
- **Shared nothing**: each computation node has its own storage.

Snowflake: big data processing platform built on top of Aws, Azure, GC.

- Architecture: persistent data stored in shared storage and each processing node has its own local storage.
- Layers:
 - Database storage: stores compressed and encrypted data that can only be accessed by Snowflake SQL.
 - Compute layer: virtual warehouses that can be scaled both horizontally and vertically containing node clusters from a cloud provider, e.g., standard, Snowpark (for large memory tasks).
 - **Virtual warehouses**: independent MPP (massive parallel processor) clusters.
 - Cloud services (top most layer): interface to other layers, handles authentication, security, query optimization. Run Snowflake SQL “cursor.execute()”
- Python connector (snowflake.connector, snowflake.connector.pandas): run Snowflake SQL and pull push data to Snowflake.

Snowsight: UI tool built on top of Snowflake (data warehousing platform) to enhance data visualization, exploration, analysis, etc.