Data Engineering Notes
Created: 2024-10-30
Updated: 2024-10-31

References:
- https://www.coursera.org/learn/spark-hadoop-snowflake-data-engineering
- https://learn.microsoft.com/en-us/training/modules/use-apache-spark-azure-databricks/
- ChatGPT 4o mini
- Google search

Apache Hadoop: open source ecosystem of software enabling parallel processing of big data.
- Hadoop Distributed File System (HDFS): storage system.
- MapReduce: framework within Hadoop that:
    - maps (distributes) tasks across a cluster of computers that store intermediate results (key-value pair) on disk (slower than in-memory).
    - Intermediate results are grouped by keys and sent to the cluster of computers (parallel processing) responsible for reducing (performing user reduction function) and results are combined.

Apache Spark: built on top of Hadoop but stores intermediate results in-memory instead of on disk. Leverages parallelism for task completion via horizontal scaling (more nodes).
- Each worker/executor node runs a JVM (multi-threaded). Each has multiple slots based on #cores #cpus of node.
- Parallelized jobs are broken down into stages to be performed in order.

Azure Databricks: parallelized data processing on Apache Spark clusters.
- A notebook instance (*SparkSession* object) controls the driver node which distributes work across worker nodes.

Resilient Distributed Datasets (RDD): immutable fundamental data structure in Spark that facilitates distributed computing.
- Resilient: fault tolerant.
- Distributed: RDDs distributed across nodes in Spark clusters.
- Operations:
    - Transformation (lazy, triggered with action): creates new RDDs after applying transformation e.g., filter.
        - Shuffle is a transformation operation that writes data to disk then re-partitions data. Shuffle defines stage boundaries.
    - Action: return result from computation, e.g., count.
- In pyspark shell (sc: spark context, spark: spark session):
    - rdd = sc.parallelize(list(range(15))) # create RDD
    - rdd2 = rdd.map(lambda x: x*2) #transformation
    - rdd3 = rdd.filter(lambda x: x%3==0) #lazy transformation
    - rdd3.count() #action triggers computation

Spark SQL DataFrames: DataSet with named columns.
- Distributed collection of data.
- Abstraction of RDDs.
- Optimized for structured data manipulation.
- Works with (semi-)structured data: CSV, Apache Avro, ORC, Parquet, JSON, relational DBs.
- PySpark integrates with Spark SQL can be used for ETL (Extract, Transform, Load) pipelines.
- Spark SQL is a module within Spark providing SQL capabilities for structured data.

- PySpark is a library allowing Spark to be used in Python.